

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
```

```
In [2]: pd.read_csv(r"C:\Users\HP\Documents\NVIDIA_STOCK.csv")
```

Out[2]:

		Price	Adj Close	Close	High	Low	Open	Volume
0	Ticker		NVDA	NVDA	NVDA	NVDA	NVDA	NVDA
1	Date		NaN	NaN	NaN	NaN	NaN	NaN
2	2018-01-02	4.929879665374756	4.983749866485596	4.987500190734863	4.862500190734863	4.894499778747559		355616000
3	2018-01-03	5.254334926605225	5.3117499351501465	5.34250020980835		5.09375	5.102499961853027	914704000
4	2018-01-04	5.2820329666137695	5.339749813079834	5.451250076293945	5.317249774932861	5.394000053405762		583268000
...	...		...	...	...	...	...	...
1694	2024-09-24	120.8616714477539	120.87000274658203	121.80000305175781	115.37999725341797	116.5199966430664		354966800
1695	2024-09-25	123.50149536132812	123.51000213623047	124.94000244140625	121.61000061035156	122.0199966430664		284692900
1696	2024-09-26	124.03145599365234	124.04000091552734	127.66999816894531	121.80000305175781	126.80000305175781		302582900
1697	2024-09-27	121.39163970947266	121.4000015258789	124.02999877929688	119.26000213623047	123.97000122070312		271009200
1698	2024-09-30	121.43163299560547	121.44000244140625		121.5	118.1500015258789	118.30999755859375	226553700

1699 rows × 7 columns

```
In [3]: NVIDIA_df= pd.read_csv(r"C:\Users\HP\Documents\NVIDIA_STOCK.csv")
```

```
In [4]: NVIDIA_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1699 entries, 0 to 1698
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  ---
 0   Price       1699 non-null   object
 1   Adj Close   1698 non-null   object
 2   Close       1698 non-null   object
 3   High        1698 non-null   object
 4   Low         1698 non-null   object
 5   Open        1698 non-null   object
 6   Volume      1698 non-null   object
dtypes: object(7)
memory usage: 93.0+ KB

```

```
In [5]: NVIDIA_df.shape
```

```
Out[5]: (1699, 7)
```

```
In [6]: NVIDIA_df.isnull().sum()
```

```

Out[6]: Price      0
Adj Close    1
Close        1
High         1
Low          1
Open         1
Volume       1
dtype: int64

```

```

In [7]: correct_columns = ['Date', 'Adj Close', 'Close', 'High', 'Low', 'Open', 'Volume']
NVIDIA_df = NVIDIA_df.iloc[2:].copy() # Skip the first two junk rows
NVIDIA_df.columns = correct_columns # Assign correct column names

```

```
In [8]: NVIDIA_df.columns.tolist()
```

```
Out[8]: ['Date', 'Adj Close', 'Close', 'High', 'Low', 'Open', 'Volume']
```

```
In [9]: # Reset index
NVIDIA_df.reset_index(drop=True, inplace=True)
```

```
In [10]: print(NVIDIA_df)
```

	Date	Adj Close	Close	High \
0	2018-01-02	4.929879665374756	4.983749866485596	4.987500190734863
1	2018-01-03	5.254334926605225	5.3117499351501465	5.34250020980835
2	2018-01-04	5.2820329666137695	5.339749813079834	5.451250076293945
3	2018-01-05	5.326793670654297	5.385000228881836	5.422749996185303
4	2018-01-08	5.490012168884277	5.550000190734863	5.625
...	...	...	...	...
1692	2024-09-24	120.8616714477539	120.87000274658203	121.80000305175781
1693	2024-09-25	123.50149536132812	123.51000213623047	124.94000244140625
1694	2024-09-26	124.03145599365234	124.04000091552734	127.66999816894531
1695	2024-09-27	121.39163970947266	121.4000015258789	124.02999877929688
1696	2024-09-30	121.43163299560547	121.44000244140625	121.5

	Low	Open	Volume
0	4.862500190734863	4.894499778747559	355616000
1	5.09375	5.102499961853027	914704000
2	5.317249774932861	5.394000053405762	583268000
3	5.2769999504089355	5.354750156402588	580124000
4	5.4644999504089355	5.510000228881836	881216000
...	...	...	...
1692	115.37999725341797	116.5199966430664	354966800
1693	121.61000061035156	122.0199966430664	284692900
1694	121.80000305175781	126.80000305175781	302582900
1695	119.26000213623047	123.97000122070312	271009200
1696	118.1500015258789	118.30999755859375	226553700

[1697 rows x 7 columns]

```
In [11]: NVIDIA_df['Date'] = pd.to_datetime(NVIDIA_df['Date'], errors='coerce')
for col in correct_columns[1:]:
    NVIDIA_df[col] = pd.to_numeric(NVIDIA_df[col], errors='coerce')
```

```
In [12]: print(NVIDIA_df.head())
```

	Date	Adj Close	Close	High	Low	Open	Volume
0	2018-01-02	4.929880	4.98375	4.98750	4.86250	4.89450	355616000
1	2018-01-03	5.254335	5.31175	5.34250	5.09375	5.10250	914704000
2	2018-01-04	5.282033	5.33975	5.45125	5.31725	5.39400	583268000
3	2018-01-05	5.326794	5.38500	5.42275	5.27700	5.35475	580124000
4	2018-01-08	5.490012	5.55000	5.62500	5.46450	5.51000	881216000

```
In [13]: NVIDIA_df.corr(numeric_only=True)
```

```
Out[13]:
```

	Adj Close	Close	High	Low	Open	Volume
Adj Close	1.000000	1.000000	0.999662	0.999763	0.999329	-0.113866
Close	1.000000	1.000000	0.999662	0.999763	0.999330	-0.113864
High	0.999662	0.999662	1.000000	0.999673	0.999780	-0.109508
Low	0.999763	0.999763	0.999673	1.000000	0.999668	-0.117262
Open	0.999329	0.999330	0.999780	0.999668	1.000000	-0.112821
Volume	-0.113866	-0.113864	-0.109508	-0.117262	-0.112821	1.000000

```
In [15]: import seaborn as sns
```

```
In [16]: sns.set()
plt.figure(figsize=(10,6))
sns.lineplot(x='Date', y='Adj Close', data=NVIDIA_df)
plt.title('Adjusted Close Price Over Time')
plt.show()
```

C:\Users\HP\anaconda3\Lib\site-packages\seaborn\\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

with pd.option\_context('mode.use\_inf\_as\_na', True):


C:\Users\HP\anaconda3\Lib\site-packages\seaborn\\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

with pd.option\_context('mode.use\_inf\_as\_na', True):



```
In [17]: plt.figure(figsize=(10,6))
sns.lineplot(x='Date', y='Volume', data=NVIDIA_df)
plt.title('Trading Volume Over Time')
plt.show()
```

```
C:\Users\HP\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
    with pd.option_context('mode.use_inf_as_na', True):  
C:\Users\HP\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
    with pd.option_context('mode.use_inf_as_na', True):
```



```
In [18]: sns.set()  
plt.figure(figsize=(10,6))  
sns.histplot(NVIDIA_df['Adj Close'], bins=50)  
plt.title('Distribution of Adjusted Close Price')  
plt.show()
```


C:\Users\HP\anaconda3\Lib\site-packages\seaborn\\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
with pd.option\_context('mode.use\_inf\_as\_na', True):



```
In [19]: plt.figure(figsize=(10,6))
sns.histplot(NVIDIA_df['Volume'], bins=50)
plt.title('Distribution of Trading Volume')
plt.show()
```



```
C:\Users\HP\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
with pd.option_context('mode.use_inf_as_na', True):
```




```
In [20]: corr_matrix = NVIDIA_df.corr()  
print(corr_matrix)
```

	Date	Adj Close	Close	High	Low	Open	\
Date	1.000000	0.766104	0.765909	0.765275	0.766219	0.765174	
Adj Close	0.766104	1.000000	1.000000	0.999662	0.999763	0.999329	
Close	0.765909	1.000000	1.000000	0.999662	0.999763	0.999330	
High	0.765275	0.999662	0.999662	1.000000	0.999673	0.999780	
Low	0.766219	0.999763	0.999763	0.999673	1.000000	0.999668	
Open	0.765174	0.999329	0.999330	0.999780	0.999668	1.000000	
Volume	-0.105869	-0.113866	-0.113864	-0.109508	-0.117262	-0.112821	

	Volume
Date	-0.105869
Adj Close	-0.113866
Close	-0.113864
High	-0.109508
Low	-0.117262
Open	-0.112821
Volume	1.000000

```
In [21]: plt.figure(figsize=(10,8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()
```




```
In [22]: plt.figure(figsize=(10,6))
sns.lineplot(x=NVIDIA_df.index, y=NVIDIA_df['Adj Close'])
plt.title('Adjusted Close Price Over Time')
plt.show()
```

C:\Users\HP\anaconda3\Lib\site-packages\seaborn\\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

with pd.option\_context('mode.use\_inf\_as\_na', True):


C:\Users\HP\anaconda3\Lib\site-packages\seaborn\\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.

with pd.option\_context('mode.use\_inf\_as\_na', True):



```
In [23]: plt.figure(figsize=(10,6))
sns.lineplot(x=NVIDIA_df.index, y=NVIDIA_df['Volume'])
plt.title('Trading Volume Over Time')
plt.show()
```

```
C:\Users\HP\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
    with pd.option_context('mode.use_inf_as_na', True):  
C:\Users\HP\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning: use_inf_as_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
    with pd.option_context('mode.use_inf_as_na', True):
```



```
In [24]: from sklearn.model_selection import train_test_split  
        from sklearn.preprocessing import StandardScaler
```

```
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

```
In [25]: target = "Close"
features = ["Adj Close", "High", "Low", "Open", "Volume"]
X = NVIDIA_df[features]
y = NVIDIA_df[target]
```

```
In [26]: # Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, shuffle=False)
```

```
In [27]: # Train Random Forest Classifier
model = RandomForestRegressor(n_estimators=100, random_state=42)
model.fit(X_train, y_train)
```

```
Out[27]: ▼      RandomForestRegressor
RandomForestRegressor(random_state=42)
```

```
In [28]: y_pred = model.predict(X_test)
```

```
In [29]: print("MSE:", mean_squared_error(y_test, y_pred))
print("MAE:", mean_absolute_error(y_test, y_pred))
print("R^2 Score:", r2_score(y_test, y_pred))
```

MSE: 2664.382655144837

MAE: 40.819758959377516

R^2 Score: -1.6669254703438692

## LINEAR REGRESSION

```
In [30]: from sklearn.linear_model import LinearRegression
```

```
In [31]: target = "Close"
features = ["Adj Close", "High", "Low", "Open", "Volume"]
X = NVIDIA_df[features]
y = NVIDIA_df[target]
```

```
In [32]: # Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, shuffle=False)
```

```
In [33]: # Train Random Forest Classifier
model = LinearRegression()
model.fit(X_train, y_train)
```

```
Out[33]: ▾ LinearRegression
LinearRegression()
```

```
In [34]: y_pred = model.predict(X_test)
```

```
In [35]: print("MSE:", mean_squared_error(y_test, y_pred))
print("MAE:", mean_absolute_error(y_test, y_pred))
print("R^2 Score:", r2_score(y_test, y_pred))
```

```
MSE: 0.00010964789972121377
MAE: 0.008906863779295674
R^2 Score: 0.999998902474553
```

```
In [36]: # Predict the entire dataset
y_pred_entire_dataset = model.predict(X)
```

```
In [37]: # Add predictions to the original dataset
NVIDIA_df['Predicted_Close'] = y_pred_entire_dataset
```

```
In [38]: print(NVIDIA_df['Predicted_Close'])
```



0	4.965087
1	5.289988
2	5.317984
3	5.362332
4	5.526071

...

1692	120.861345
1693	123.512843
1694	124.044186
1695	121.403287
1696	121.436479

Name: Predicted\_Close, Length: 1697, dtype: float64

In [ ]: